



南京大學

NANJING UNIVERSITY

# 运输层

殷亚凤

智能软件与工程学院

苏州校区南雍楼东区225

yafeng@nju.edu.cn , <https://yafengnju.github.io/>



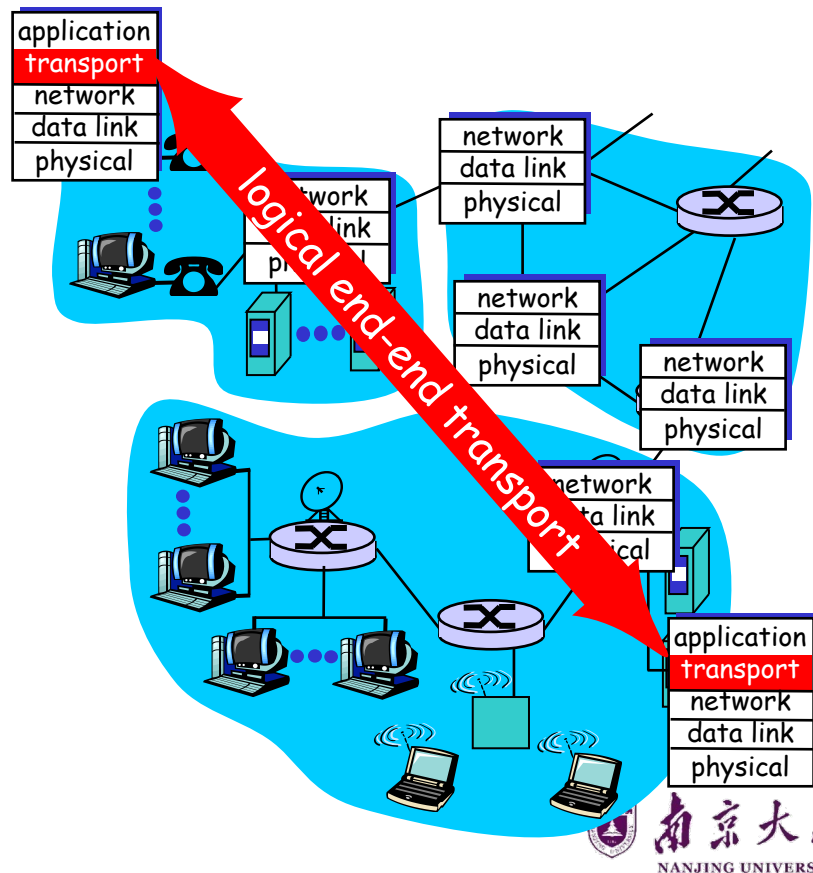
- Transport layer basics
- Design of reliable transport
- Designing a reliable transport protocol





# Internet Transport Services

- Provide **logical communication** between app processes running on different hosts
- Transport protocols run in **end systems**
  - **Send side**: breaks app messages into segments, passes to network layer
  - **Receive side**: reassembles segments into messages, passes to app layer
- More than one transport protocol available to apps
  - Internet: TCP and UDP





# Why a transport layer?

- IP packets are addressed to a **host** but **end-to-end communication** is between **application processes** at hosts
  - Need a way to decide which packets go to which applications (multiplexing/demultiplexing)
- IP provides a **weak service model** (best-effort)
  - Packets can be corrupted, delayed, dropped, reordered, duplicated
  - No guidance on how much traffic to send and when
  - Dealing with this is tedious for application developers





# Multiplexing & demultiplexing

---

- Multiplexing (Mux)
  - Gather and combining data chunks at the source host from different applications and delivering to the network layer
- Demultiplexing (Demux)
  - Delivering correct data to corresponding sockets from multiplexed a stream



# Role of the transport layer

---

- Communication between processes
  - Mux and demux from/to application processes
  - Implemented using ports





# Role of the transport layer

---

- Communication between processes
- Provide common end-to-end services for app layer [optional]
  - Reliable, in-order data delivery
  - Well-paced data delivery
    - Too fast may overwhelm the network
    - Too slow is not efficient





# Role of the transport layer

---

- Communication between processes
- Provide common end-to-end services for app layer [optional]
- TCP and UDP are the common transport protocols
  - Also SCTP, MPTCP, SST, RDP, DCCP, ...







# Role of the transport layer

---

- Communication between processes
- Provide common end-to-end services for app layer [optional]
- TCP and UDP are the common transport protocols
- **UDP is a minimalist transport protocol**
  - Only provides mux/demux capabilities





# Role of the transport layer

---

- Communication between processes
- Provide common end-to-end services for app layer [optional]
- TCP and UDP are the common transport protocols
- UDP is a minimalist transport protocol
- **TCP offers a reliable, in-order, byte stream abstraction**
  - With congestion control, but w/o performance guarantees (delay, b/w, etc.)





# Applications and sockets

---

- **Socket**: **software abstraction** for an **application process** to exchange network messages with the (transport layer in the) **operating system**
- Transport layer addressing
  - $\langle \text{HostIP}, \text{Port} \rangle$ , called a socket
- Two important types of sockets
  - UDP socket: TYPE is `SOCK_DGRAM`
  - TCP socket: TYPE is `SOCK_STREAM`





# Ports

- 16-bit numbers that help distinguishing apps
  - Packets carry **src/dst port** no in transport header
  - Well-known (0-1023) and ephemeral ports
- OS stores **mapping** between **sockets and ports**
  - Port in packets and sockets in OS
  - For **UDP ports (SOCK\_DGRAM)**
    - OS stores (local port, local IP address)  $\leftrightarrow$  socket
  - For **TCP ports (SOCK\_STREAM)**
    - OS stores (local port, local IP, remote port, remote IP)  $\leftrightarrow$  socket





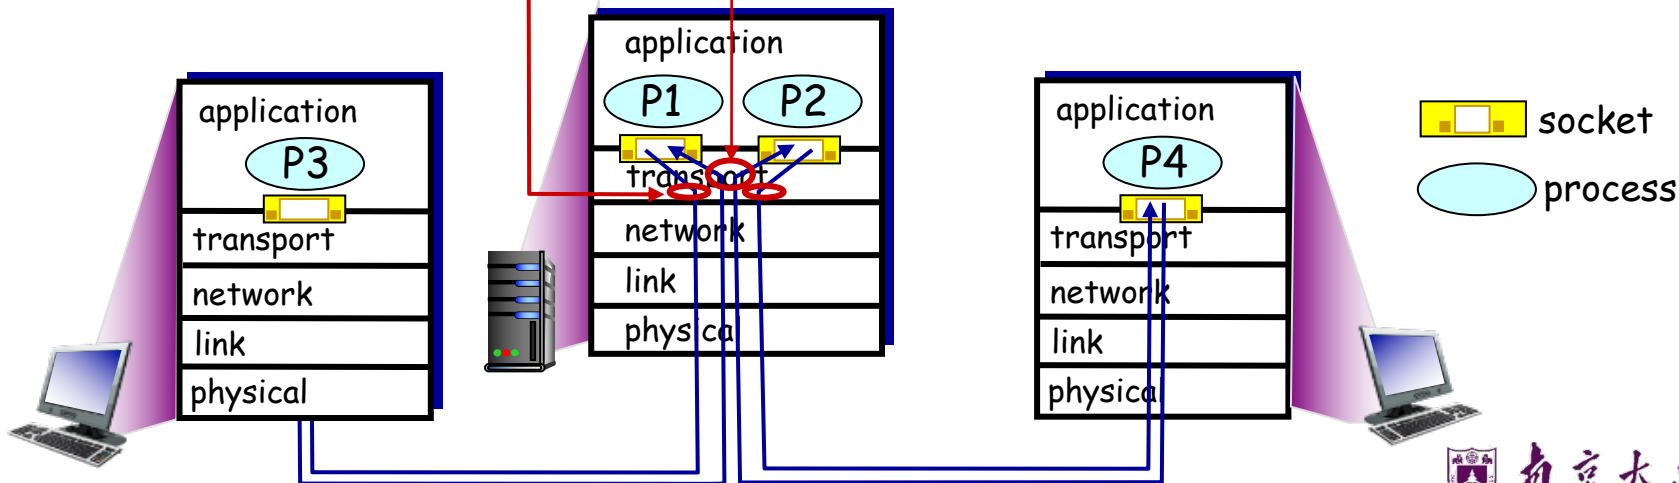
# Multiplexing/demultiplexing

## multiplexing at sender:

handle data from multiple sockets, add transport header (later used for demultiplexing)

## demultiplexing at receiver:

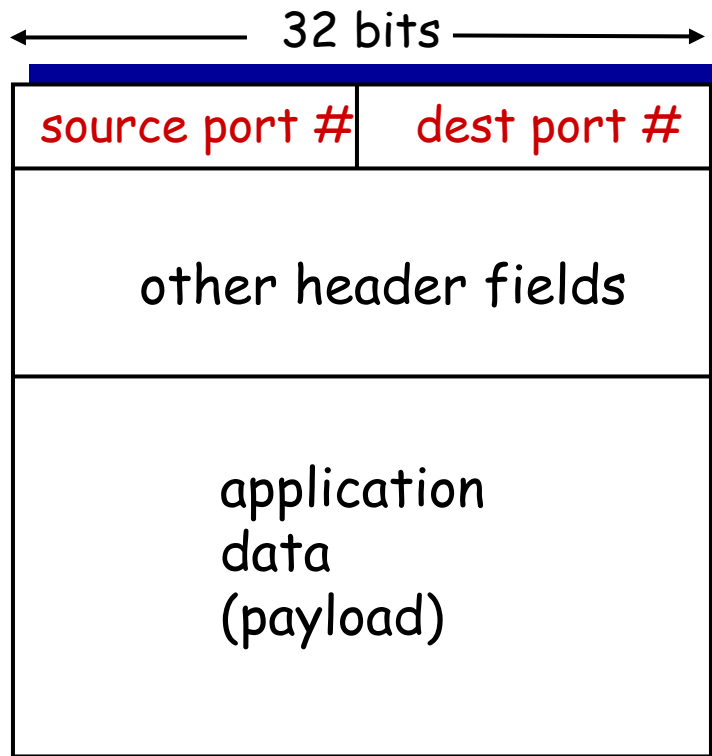
use header info to deliver received segments to correct socket





# How demultiplexing works

- host receives **IP datagrams**
  - each datagram has **source IP address, destination IP address**
  - each datagram carries one **transport-layer segment**
  - each segment has source, destination **port number**
- host uses **IP addresses & port numbers** to direct segment to appropriate socket



TCP/UDP segment format 大學



# Connectionless demultiplexing

- **recall**: created socket has host-local port #:  

```
DatagramSocket mySocket1  
= new DatagramSocket(12534);
```
- **recall**: when creating datagram to send into UDP socket, must specify
  - destination IP address
  - destination port #
- When host receives UDP segment:
  - checks destination port # in segment
  - directs UDP segment to socket with that port #



IP datagrams with **same dest. port #**, but different source IP addresses and/or source port numbers will be directed to **same socket** at dest.



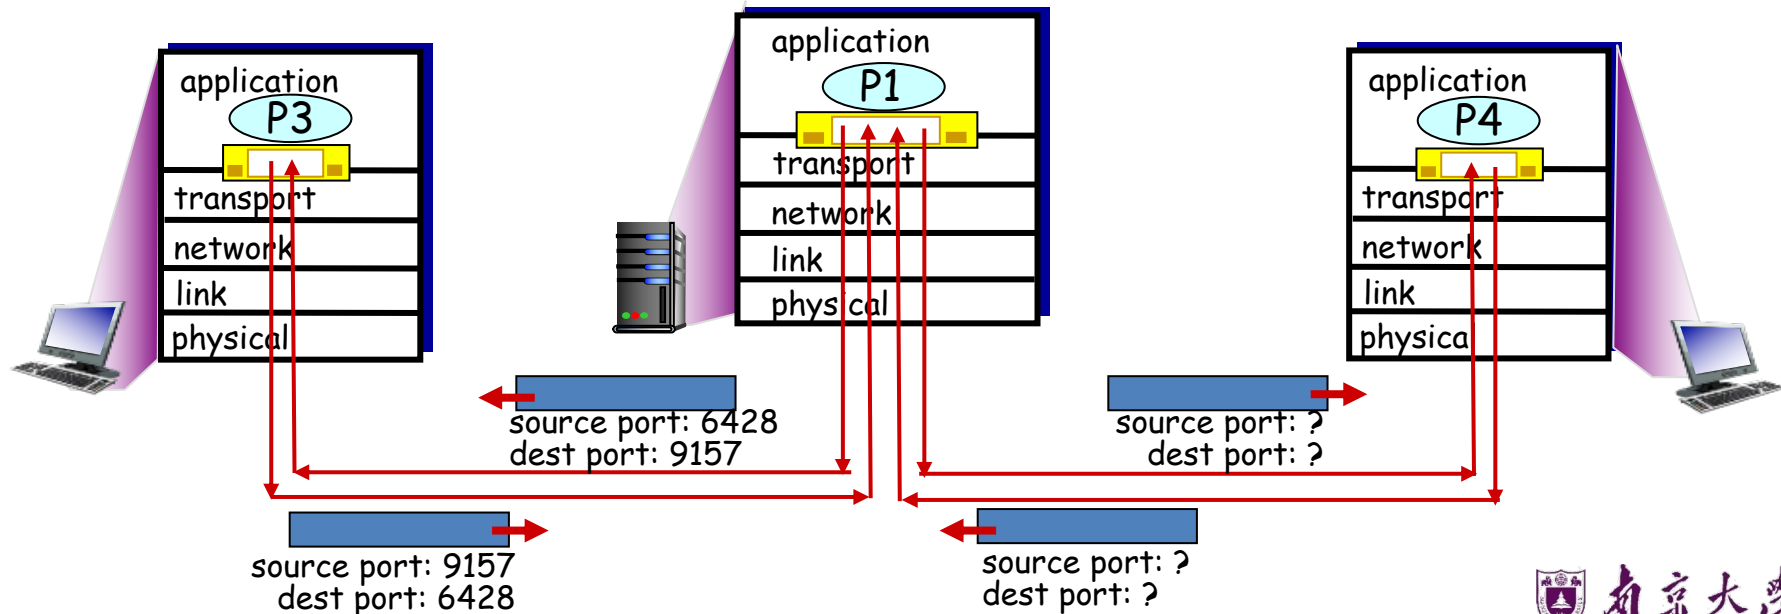


# Connectionless demux: example

```
DatagramSocket  
mySocket2 = new  
DatagramSocket  
(9157);
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428);
```

```
DatagramSocket  
mySocket1 = new  
DatagramSocket  
(5775);
```







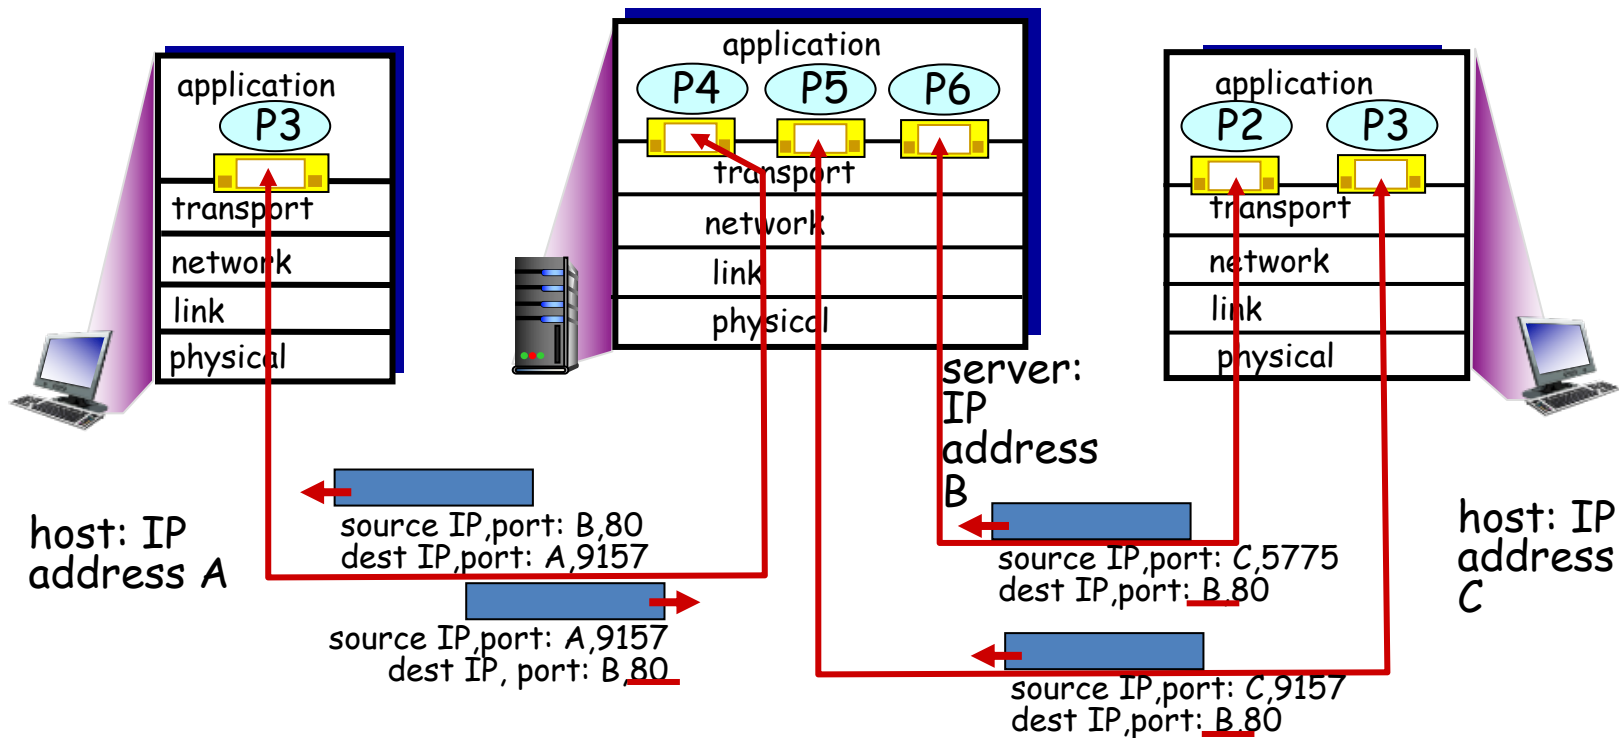
# Connection-oriented demux

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- demux: receiver **uses all four values** to direct segment to appropriate socket
- server host may support many **simultaneous TCP sockets**:
  - each socket identified by its own 4-tuple
- web servers have different sockets for each connecting client
  - non-persistent HTTP will **have different socket for each request**





# Connection-oriented demux: example

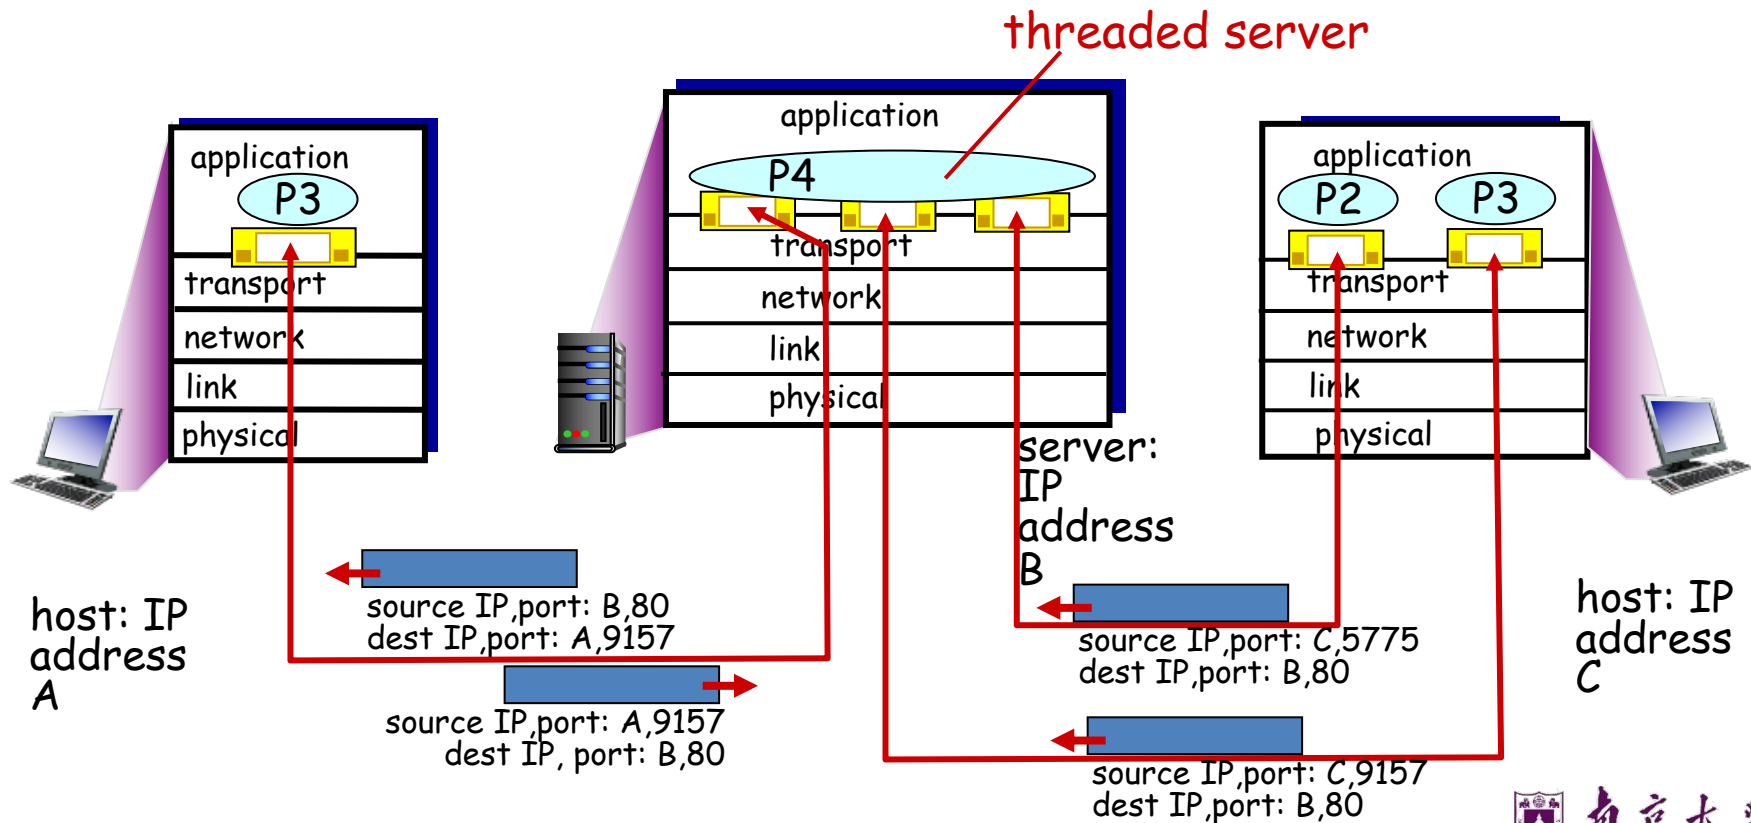


three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets





# Connection-oriented demux: example





- Transport layer basics
- Design of reliable transport
- Designing a reliable transport protocol





# Reliable transport

- In a perfect world, reliable transport is easy

@Sender

– Send packets

@Receiver

– Wait for packets





# Reliable transport

---

- All the bad things best-effort can do
  - A packet is **corrupted** (bit errors)
  - A packet is **lost** (why?)
  - A packet is **delayed** (why?)
  - Packets are **reordered** (why?)
  - A packet is **duplicated** (why?)





# Reliable transport

---


- Mechanisms for coping with bad events
  - **Checksums**: to detect corruption
  - **ACKs**: receiver tells sender that it received packet
  - **NACK**: receiver tells sender it did not receive packet
  - **Sequence numbers**: a way to identify packets
  - **Retransmissions**: sender resends packets
  - **Timeouts**: a way of deciding when to resend packets
  - **Forward error correction**: a way to mask errors without retransmission
  - **Network encoding**: an efficient way to repair errors





# Dealing with packet corruption

- the question: how to recover from errors:
  - acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
  - negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
  - sender retransmits pkt on receipt of NAK








# Dealing with packet corruption

What if the ACK/NACK is corrupted?

Data and ACK packets carry sequence numbers






# Dealing with packet loss


## Timer-driven loss detection

Set timer when packet is sent; retransmit on timeout






# Dealing with packet loss (of ack)





# Dealing with packet loss

Timer-driven retransmission can lead to duplicates





# Components of a solution

---

- Checksums (to detect bit errors)
- Timers (to detect loss)
- Acknowledgements (positive or negative)
- Sequence numbers (to deal with duplicates)





- Transport layer basics
- Design of reliable transport
- Designing a reliable transport protocol





# A Solution: "Stop and Wait"

## @Sender

- Send packet(I); (re)set timer; wait for ack
- If (ACK)
  - I++; repeat
- If (NACK or TIMEOUT)
  - repeat

## @Receiver


- Wait for packet
- If packet is OK, send ACK
- Else, send NACK
- Repeat

- A correct reliable transport protocol, but an extremely inefficient one





# Stop & Wait is inefficient



$L$ : packet size  
 $R$ : bandwidth of the link  
 $RTT = 2 * PropDelay$ : roundtrip time

If  $(L/R \ll RTT)$  then  
Throughput  $\sim DATA/RTT$





# Orders of magnitude

- e.g.: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:

$$D_{\text{trans}} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$

- if RTT=30 msec,
- $U_{\text{sender}}$ : *utilization* - fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- 33kB/sec throughput over 1 Gbps link!
- network protocol limits use of physical resources!






# Pipelined protocols

**pipelining:** sender allows multiple, “in-flight”, yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation



## Pipelining: increased utilization



$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R} = \frac{.0024}{30.008} = 0.00081$$



# Three design decisions

---

- Which packets can **sender send**?
  - Sliding window
- How does receiver **ack packets**?
  - Cumulative
  - Selective
- Which packets does **sender resend**?
  - Go-Back N (GBN)
  - Selective Repeat (SR)





# Sliding window


- **Window** = set of adjacent sequence numbers
  - The size of the set is the window size; assume window size is  $n$
- General idea: **send up to  $n$  packets at a time**
  - **Sender** can send packets in its window
  - **Receiver** can accept packets in its window
  - **Window** of acceptable packets "slides" on successful reception/acknowledgement
  - **Window** contains all packets that might still be in transit
- Sliding window often called "packets in flight"






# Sliding window

- Let  $A$  be the **last ack'd packet of sender without gap**; then window of sender =  $\{A+1, A+2, \dots, A+n\}$



- Already ACK'd
- Sent but not ACK'd
- Cannot be sent

- Let  $B$  be the **last received packet without gap** by receiver, then window of receiver =  $\{B+1, \dots, B+n\}$



- Received and ACK'd
- Acceptable but not yet received
- Cannot be received



# Throughput of sliding window

---

- If window size is  $n$ , then throughput is roughly
  - $\text{MIN}(n * \text{DATA} / \text{RTT}, \text{Link Bandwidth})$
- Compare to Stop and Wait:  $\text{Data} / \text{RTT}$
- What happens when  $n$  gets too large?





# Acknowledgements w/ sliding window

---


- Two common options
  - **Cumulative ACKs**: ACK carries next in-order sequence number that the receiver expects






# Cumulative acknowledgements

- At receiver



- Received and ACK'd
- Acceptable but not yet received
- Cannot be received

- After receiving B+1, B+2




- Receiver sends  $ACK(B+3) = ACK(B_{new}+1)$




# Cumulative acknowledgements (cont'd)

- At receiver



- Received and ACK'd
- Acceptable but not yet received
- Cannot be received

- After receiving  $B+4$ ,  $B+5$



- Receiver sends  $ACK(B+1)$



# Acknowledgements w/ sliding window

- Two common options
  - **Cumulative ACKs**: ACK carries **next in-order** sequence number the receiver expects
  - **Selective ACKs**: ACK individually acknowledges **correctly received** packets
- Selective ACKs offer more precise information but require more complicated book-keeping





# Sliding window protocols

---

- Resending packets: two canonical approaches
  - Go-Back-N
  - Selective Repeat
- Many variants that differ in implementation details





# Go-Back-N (GBN)

---

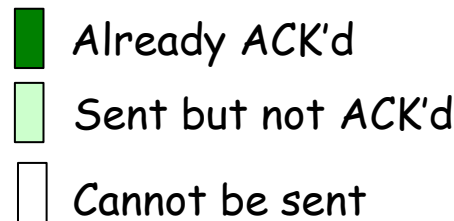

- **Sender** transmits up to  $n$  unacknowledged packets
- **Receiver** only **accepts packets in order**
  - Discards out-of-order packets (i.e., packets other than  $B+1$ )
- Receiver uses **cumulative acknowledgements**
  - i.e., sequence# in ACK = next expected in-order sequence#
- Sender sets timer for 1st outstanding ack ( $A+1$ )
- If timeout, retransmit  $A+1, \dots, A+n$



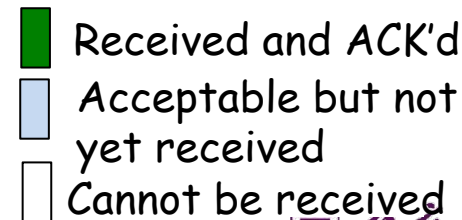



# Sliding window with GBN

- Let  $A$  be the last ack'd packet of sender without gap;  
then window of sender =  $\{A+1, A+2, \dots, A+n\}$



- Let  $B$  be the last received packet without gap by receiver, then window of receiver =  $\{B+1, \dots, B+n\}$





# GBN example w/o errors

Sender Window

$\{1\}$   
 $\{1, 2\}$   
 $\{1, 2, 3\}$   
 $\{2, 3, 4\}$   
 $\{3, 4, 5\}$   
 $\{4, 5, 6\}$   
.  
.  
.

1  
2  
3  
4  
5  
6  
.  
.  
.

Window size = 3 packets

Receiver Window

Sender


Time

Receiver





# GBN example with errors







# Selective Repeat (SR)

- **Sender**: transmit up to **n unacknowledged** packets
- Assume packet  $k$  is lost,  $k+1$  is not
  - **Receiver**: indicates **packet  $k+1$  correctly received**
  - **Sender**: **retransmit only packet  $k$**  on timeout
- Efficient in retransmissions but complex book-keeping
  - Need a timer per packet




# SR example with errors

Sender Window

Window size = 3 packets

Receiver Window





# GBN vs. Selective Repeat

---

- When would **GBN** be better?
  - When **error rate is low**; wastes bandwidth otherwise
- When would **SR** be better?
  - When **error rate is high**; otherwise, too complex



# 提问

## Q & A

殷亚凤

智能软件与工程学院

苏州校区南雍楼东区225

yafeng@nju.edu.cn , <https://yafengnju.github.io/>



南京大學  
NANJING UNIVERSITY