



南京大學

NANJING UNIVERSITY

# 运输层

殷亚凤

智能软件与工程学院

苏州校区南雍楼东区225

yafeng@nju.edu.cn , <https://yafengnju.github.io/>



# Outline

---

- UDP: User Datagram Protocol
- TCP: Transmission Control Protocol
- TCP Connection Setup
- TCP Connection Teardown





# UDP: User Datagram Protocol

---

- Lightweight **communication** between **processes**
  - Avoid overhead and delays of order & reliability
- UDP described in RFC 768 - (1980!)
  - Destination **IP address and port** to support demultiplexing





# UDP (cont'd)

---

- “Best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- Connectionless:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others
- UDP use:
  - streaming multimedia apps (loss tolerant, rate sensitive)
  - DNS
  - SNMP





# Why is there a UDP?

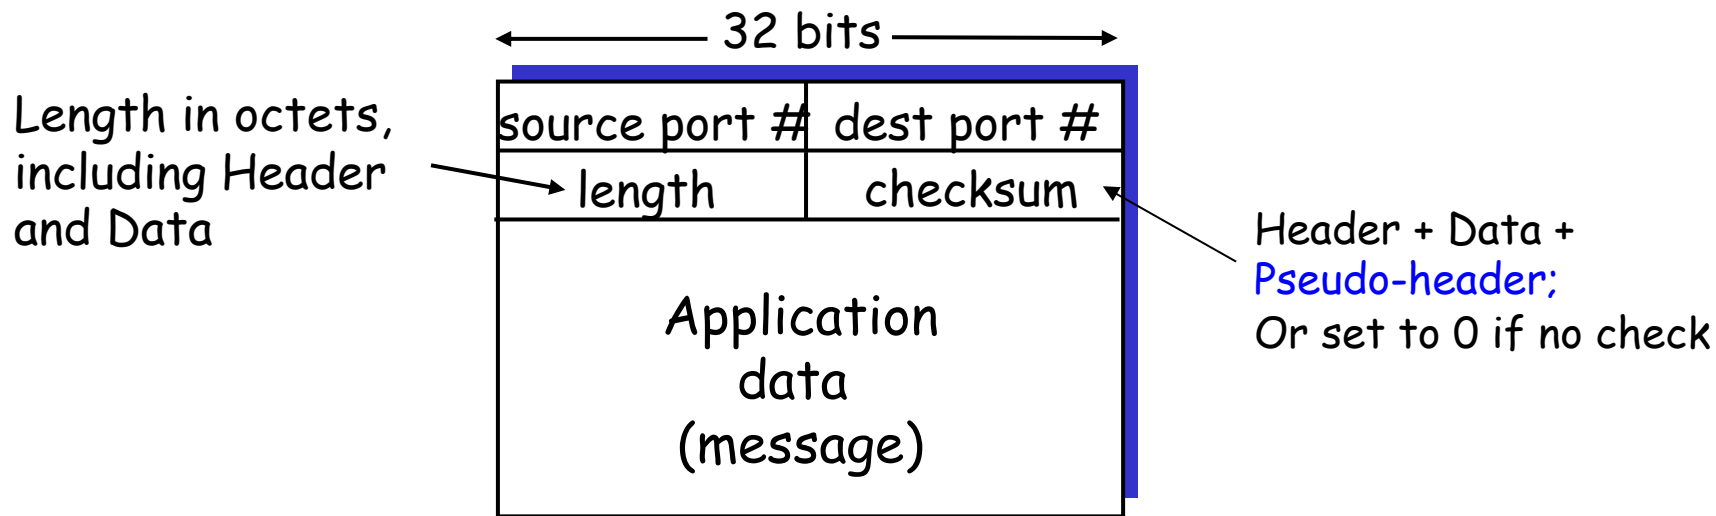
---

- no connection establishment (which can add delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control: UDP can blast away as fast as desired





# UDP Segment Format





# UDP checksum

**Goal:** detect “errors” (e.g., flipped bits) in transmitted segment

**sender:**

- treat segment contents, including header fields, as sequence of 16-bit integers
- **checksum:** addition of segment contents, and its complement sum
- sender puts checksum value into UDP checksum field

**receiver:**

- compute **checksum** of received segment
- check if the sum of computed checksum and checksum field value equals 1111....1111:
  - NO - error detected
  - YES - no error detected. **But maybe errors nonetheless?**





# Internet checksum: example

example: add two 16-bit integers

1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1

wraparound 1 1 0 1 1 1 0 1 1 1 0 1 1 1 0 1 1

sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

*Note:* when adding numbers, a carryout from the most significant bit needs to be added to the result





# Outline

---

- UDP: User Datagram Protocol
- TCP: Transmission Control Protocol
- TCP Connection Setup
- TCP Connection Teardown





# The TCP Abstraction

---

- TCP delivers a **reliable, in-order, byte stream**
- **Reliable**: TCP resends lost packets (recursively)
  - Until it gives up and shuts down connection
- **In-order**: TCP only hands consecutive chunks of data to application
- **Byte stream**: TCP assumes there is an incoming stream of data, and attempts to deliver it to app





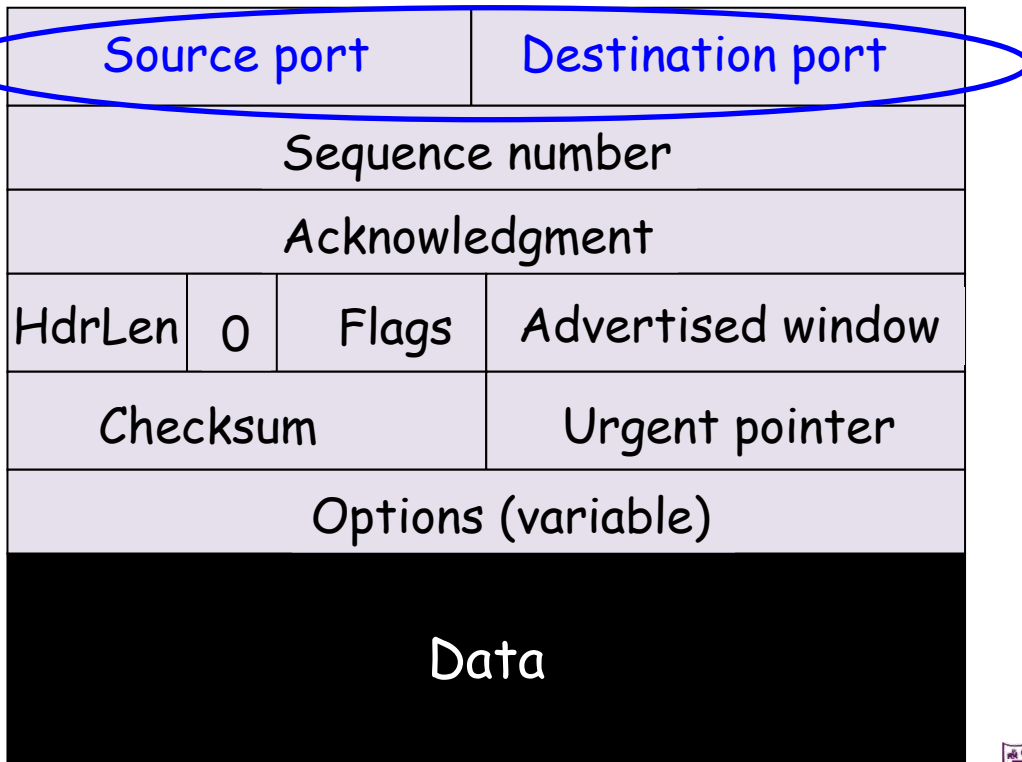
# What does TCP use from what we've seen so far?

- Most of what we've seen
  - Checksums
  - Sequence numbers are byte offsets
  - Sender and receiver maintain a sliding window
  - Receiver sends cumulative acknowledgements (like GBN)
    - ✓ Sender maintains a single retransmission timer
  - Receivers buffer out-of-sequence packets (like SR)
- Few more: fast retransmit, timeout estimation algorithms etc.



# TCP header

Used to Mux  
and Demux





# TCP header

Computed  
over pseudo-header  
and data

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			





# What does TCP do?

---

- Most of what we've seen
  - Checksum
  - Sequence numbers are byte offsets





# TCP header

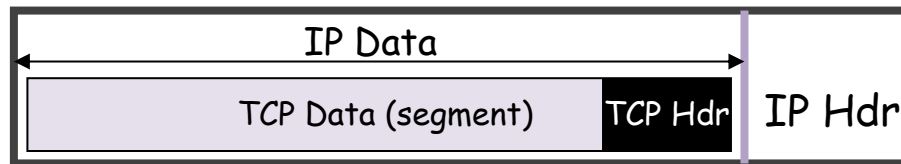
Byte offsets  
(NOT packet id),  
because TCP is a  
byte stream

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLe	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			





# TCP segment



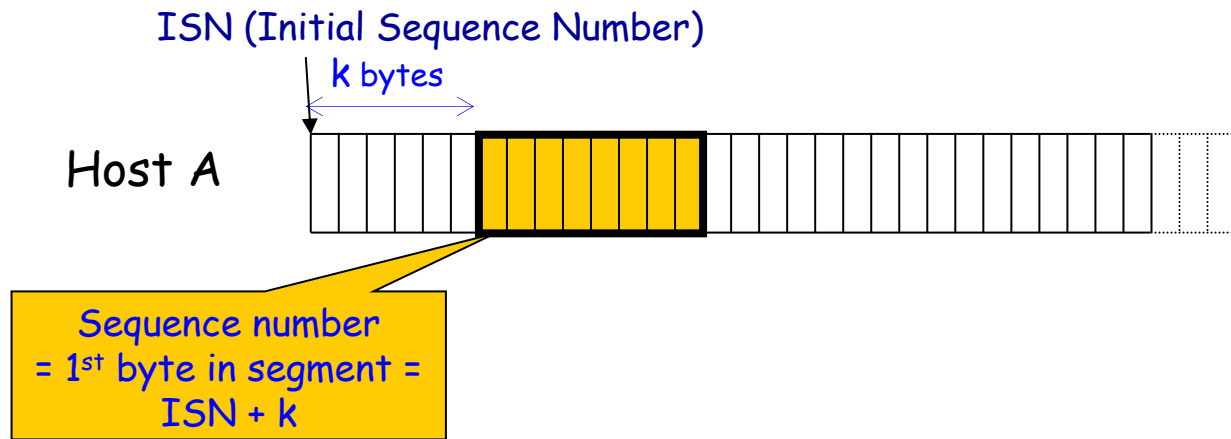
- IP packet
  - No bigger than **Maximum Transmission Unit (MTU)**
  - E.g., up to 1500 bytes with Ethernet
- TCP packet
  - IP packet with a TCP header and data inside
  - TCP header  $\geq 20$  bytes long
- TCP segment
  - No more than **Maximum Segment Size (MSS)** bytes
  - E.g., up to 1460 consecutive bytes from the stream
  - $MSS = MTU - (IP \text{ header}) - (TCP \text{ header})$





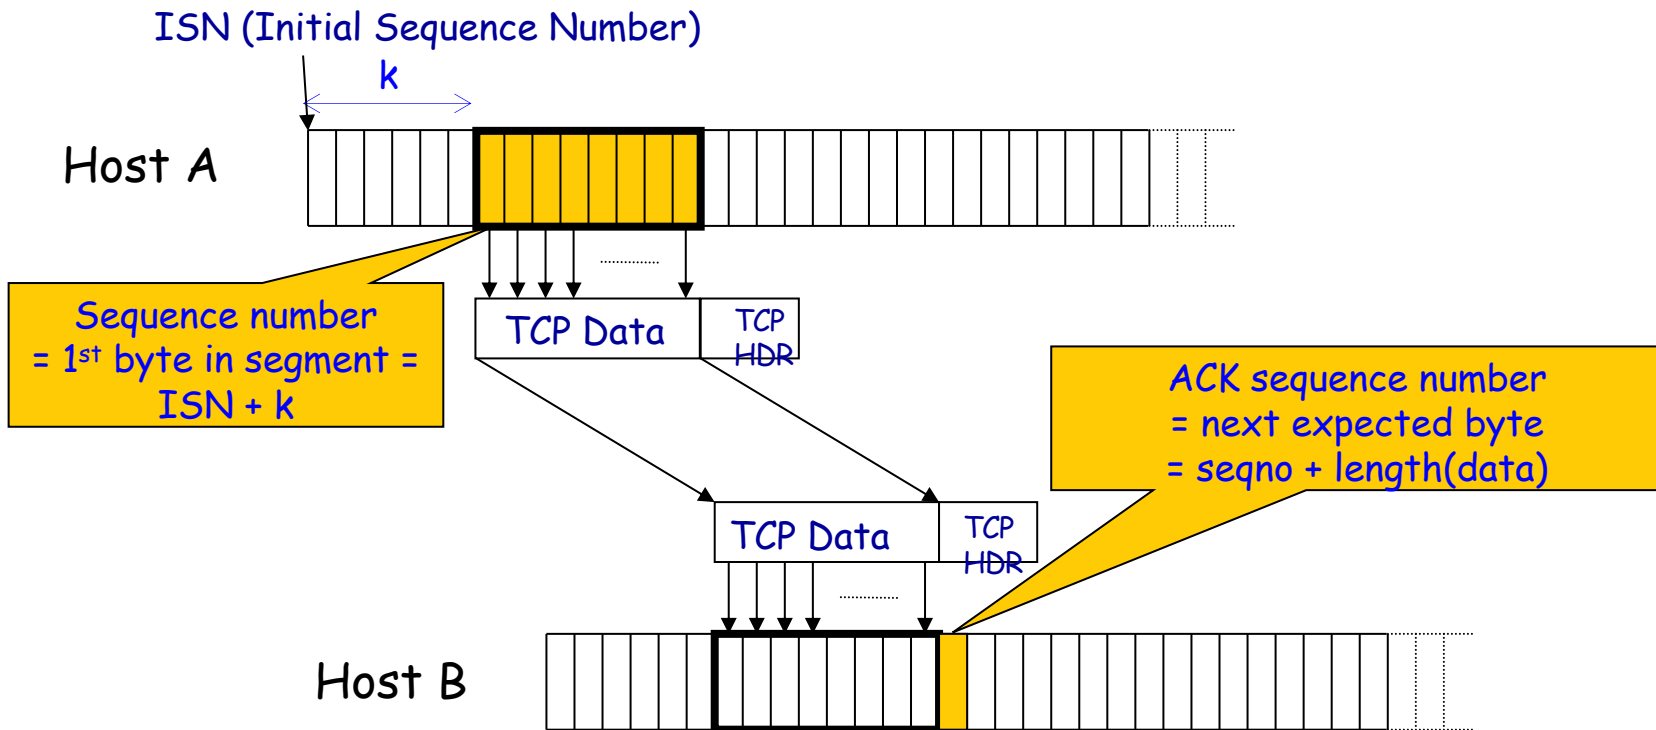


# Sequence numbers





# Sequence numbers





# What does TCP do?

---

- Most of what we've seen
  - Checksum
  - Sequence numbers are byte offsets
  - Receiver sends **cumulative acknowledgements** (like GBN)





# ACKs and sequence numbers

- **Sender sends packet**
  - Data starts with sequence number  $X$
  - Packet contains  $B$  bytes  $[X, X+1, X+2, \dots, X+B-1]$
- Upon receipt of packet, **receiver sends an ACK**
  - If all data prior to  $X$  already received:
    - ✓ ACK acknowledges  $X+B$  (because that is next expected byte)
  - If highest in-order byte received is  $Y$  s.t.  $(Y+1) < X$ 
    - ✓ ACK acknowledges  $Y+1$
    - ✓ Even if this has been ACKed before



# Typical operation

- **Sender:**  $\text{seqno} = X$ ,  $\text{length} = B$
- **Receiver:**  $\text{ACK} = X + B$
- **Sender:**  $\text{seqno} = X + B$ ,  $\text{length} = B$
- **Receiver:**  $\text{ACK} = X + 2B$
- **Sender:**  $\text{seqno} = X + 2B$ ,  $\text{length} = B$
- **Seqno of next packet is same as last ACK field**





# TCP header

Acknowledgment  
gives seqno just  
beyond highest  
seqno received in  
order

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			





# What does TCP do?

---

- Most of what we've seen
  - Checksum
  - Sequence numbers are byte offsets
  - Receiver sends cumulative acknowledgements (like GBN)
  - Receivers **can buffer out-of-sequence packets** (like SR)





# What does TCP introduce?

---

- Most of what we've seen
  - Checksum
  - Sequence numbers are byte offsets
  - Receiver sends cumulative acknowledgements (like GBN)
  - Receivers can buffer out-of-sequence packets (like SR)
- Introduces **fast retransmit**: duplicate ACKs trigger early retransmission





# Loss with cumulative ACKs

- Duplicate ACKs are a sign of an isolated loss
  - The lack of ACK progress means 500 hasn't been delivered
  - Stream of ACKs means some packets are being delivered
- Trigger retransmission upon receiving  $k$  duplicate ACKs
  - TCP uses  $k=3$
  - Faster than waiting for timeout



# Loss with cumulative ACKs

- Two choices after **resending**:
  - **Send missing packet and move sliding window** by the number of dup ACKs
    - ✓ Speeds up transmission, but might be wrong
  - **Send missing packet, and wait for ACK** to move sliding window
    - ✓ Is slowed down by single dropped packets
- Which should TCP do?





# What does TCP introduce?

---

- Most of what we've seen
  - Checksum
  - Sequence numbers are byte offsets
  - Receiver sends cumulative acknowledgements (like GBN)
  - Receivers buffer out-of-sequence packets (like SR)
- Introduces fast retransmit: duplicate ACKs trigger early retransmission
- Sender maintains a **single retransmission timer** (like GBN) and retransmits on timeout





# Retransmission timeout

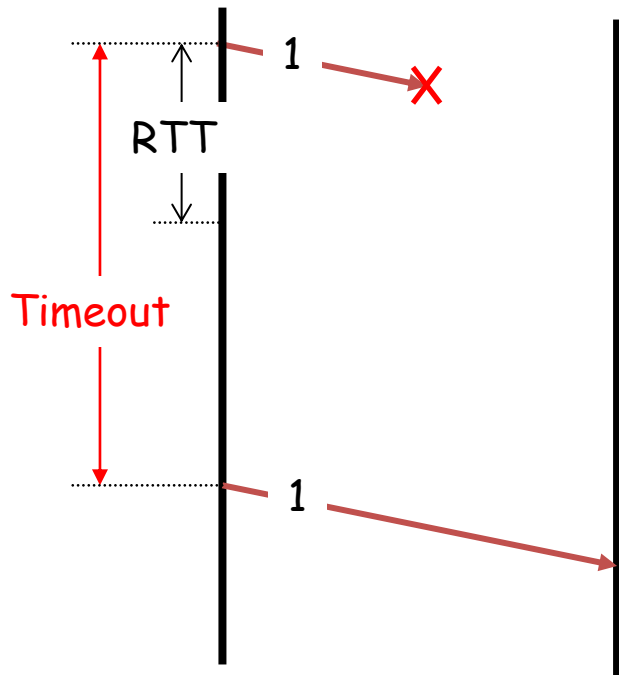
---

- If the sender hasn't received an ACK by timeout, **retransmit the first packet** in the window
- How do we pick a **timeout value**?

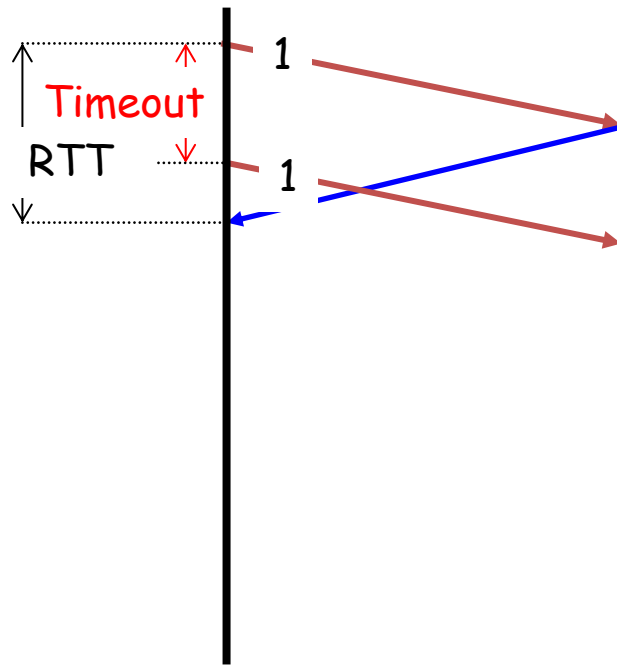




# Timing illustration



Timeout too long → inefficient



Timeout too short → duplicate packets



# Retransmission timeout

---

- If the sender hasn't received an ACK by timeout, retransmit the first packet in the window
- How to set timeout?
  - Too long: connection has low throughput
  - Too short: retransmit packet that was just delayed
- Solution: make timeout proportional to RTT
  - But how do we measure RTT?

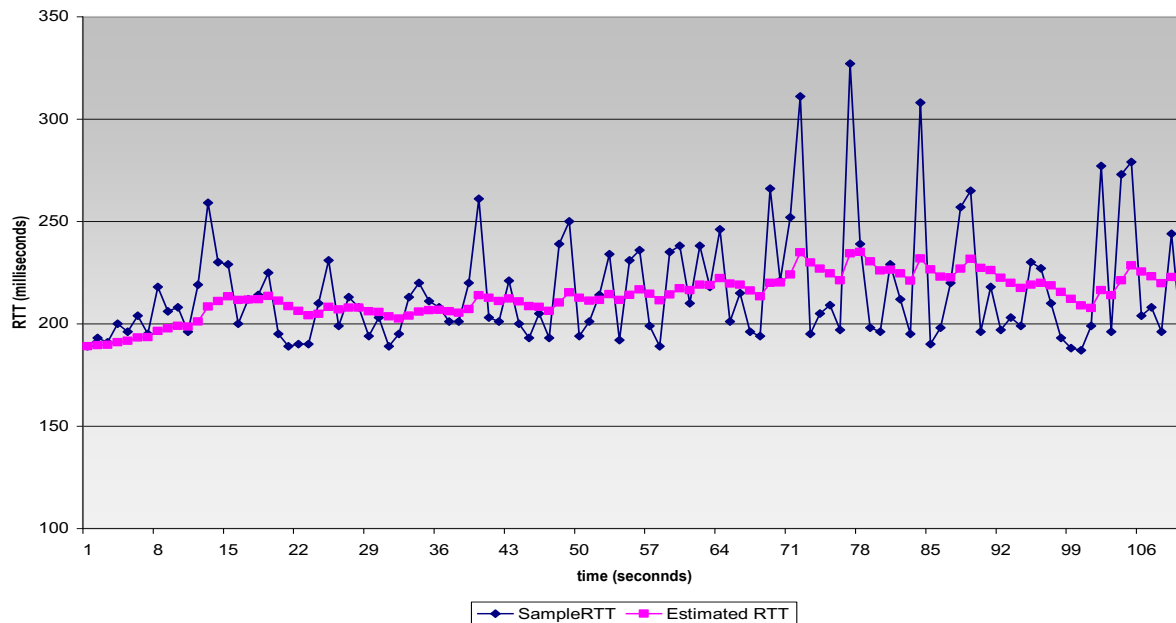




# RTT estimation

- Exponential weighted average of RTT samples

$$\text{EstimatedRTT} = (1 - \alpha) * \text{EstimatedRTT} + \alpha * \text{SampleRTT}$$





# Jacobson/Karels algorithm

- **Problem**: need to better capture variability in RTT
  - Directly measure deviation
- **Deviation** =  $| \text{SampleRTT} - \text{EstimatedRTT} |$
- **DevRTT**: exponential average of Deviation
- **RTO** =  $\text{EstimatedRTT} + 4 \times \text{DevRTT}$

$$SRTT(k+1) = (1-g) \times SRTT(k) + g \times RTT(k+1)$$

$$SERR(k+1) = RTT(k+1) - SRTT(k)$$

$$SDEV(k+1) = (1-h) \times SDEV(k) + h \times |SERR(k+1)|$$

$$RTO(k+1) = SRTT(k+1) + f \times SDEV(k+1)$$

$$g = \frac{1}{8} = 0.125 \quad h = \frac{1}{4} = 0.25 \quad f = 2 \text{ or } 4$$







# Outline

---

- UDP: User Datagram Protocol
- TCP: Transmission Control Protocol
- TCP Connection Setup
- TCP Connection Teardown





- TCP header field for connection establishment and teardown

Source Port					Destination Port				
Sequence Number									
Acknowledgement Number									
Data Offset	Reserved	U	A	P	R	S	F	Window	
		R	C	S	S	Y	I		
		G	K	H	T	N	N		
Checksum					Urgent Pointer				
TCP Options								Padding	
Data									



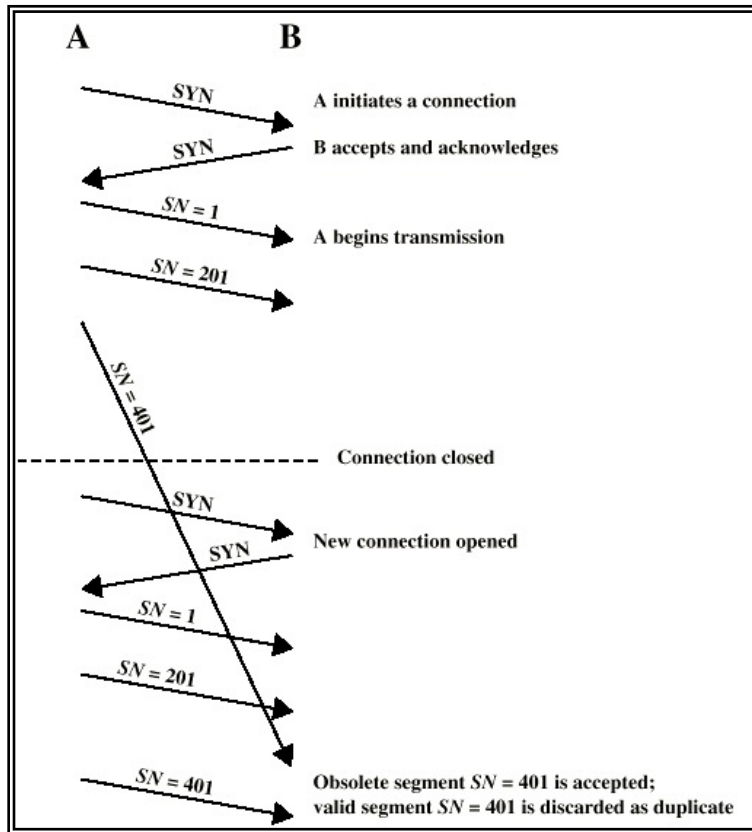
# Connection Establishment

- 2-way handshake
  - A sends SYN, B replies with SYN
  - Lost SYNs handled by re-transmission
  - Ignore duplicate SYNs once connected
- Problem
  - How to recognize slipped segments from old connection
  - How to recognize duplicated obsolete SYN





# 2-Way Handshake: Slipped Data Segment





# Initial Sequence Number (ISN)

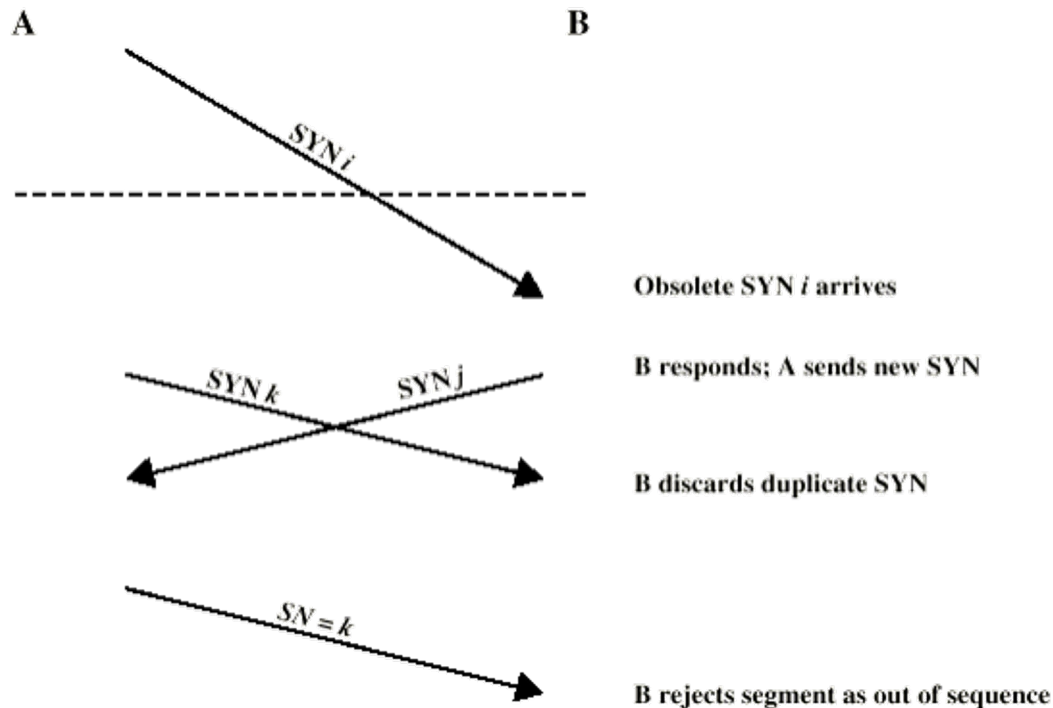
---

- Handle
  - Start each new connection with a different **initial sequence number (ISN)** far from previous connection
  - The connection request is of the form SYN  $i+1$ , where  $i$  is the sequence number of the first data segment that will be sent on this connection.
- However:





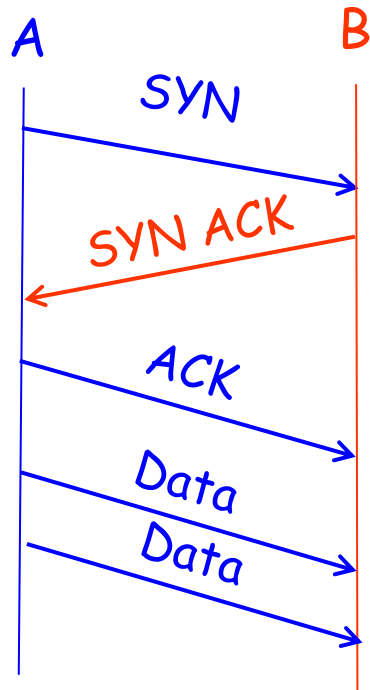
# 2-Way Handshake: Obsolete SYN





# Solution: three-way handshake

- Three-way handshake to establish connection
  - Host A sends a SYN (open; “synchronize sequence numbers”) to host B
  - Host B returns a SYN acknowledgment (SYN ACK)
  - Host A sends an ACK to acknowledge the SYN ACK



三方握手：确认对方的SYN和序号



# TCP header

Flags:

SYN

ACK

FIN

RST

PSH

URG

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			







# Step 1: A's initial SYN packet

A tells B to open  
a connection

A's port			B's port
A's Initial Sequence Number			
N/A			
5	0	SYN	Advertised window
Checksum			Urgent pointer





# Step 1: B's SYN-ACK packet

B tells it accepts  
and is ready to  
accept next  
packet

B's port		A's port	
B's Initial Sequence Number			
ACK=A's ISN+1			
5	0	SYN ACK	Advertised window
Checksum		Urgent pointer	





# Step 1: A's ACK to SYN-ACK

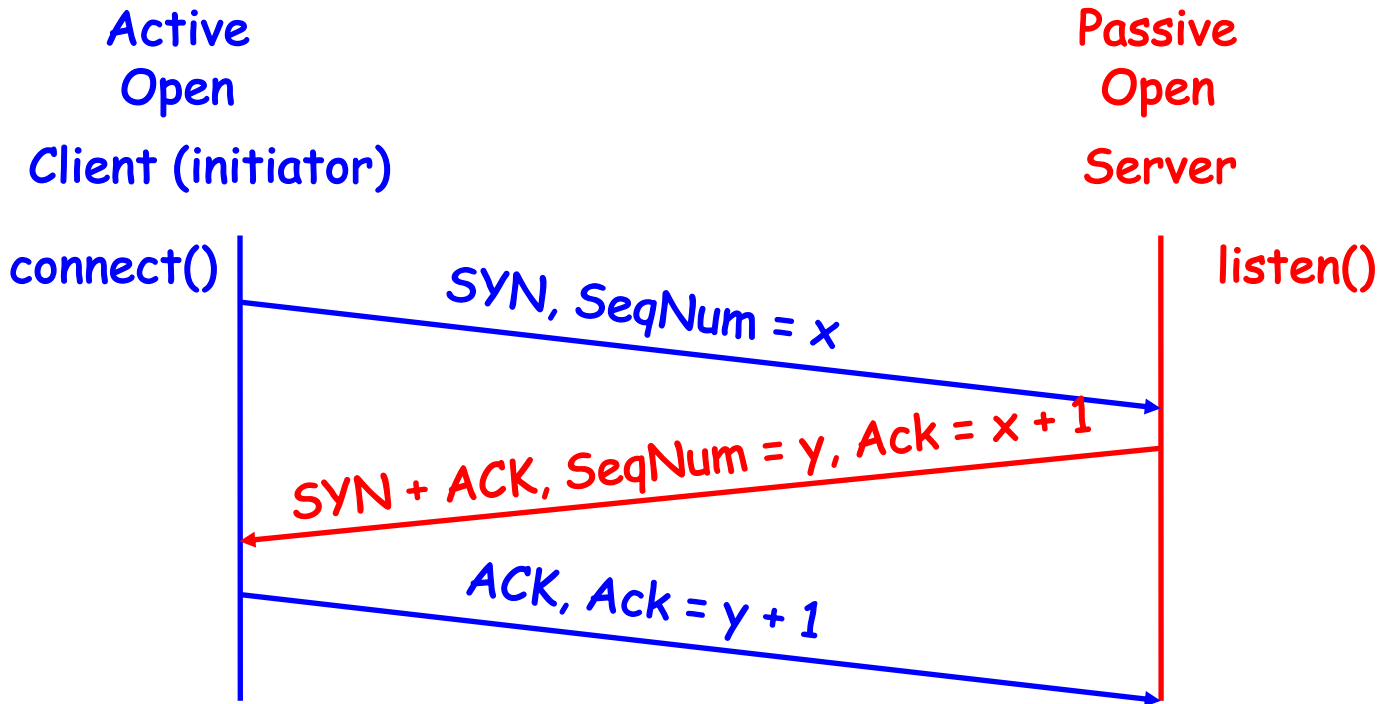
A tells B to open  
a connection

A's port			B's port
A's Initial Sequence Number			
ACK=B's ISN+1			
5	0	ACK	Advertised window
Checksum			Urgent pointer





# TCP's 3-Way handshaking





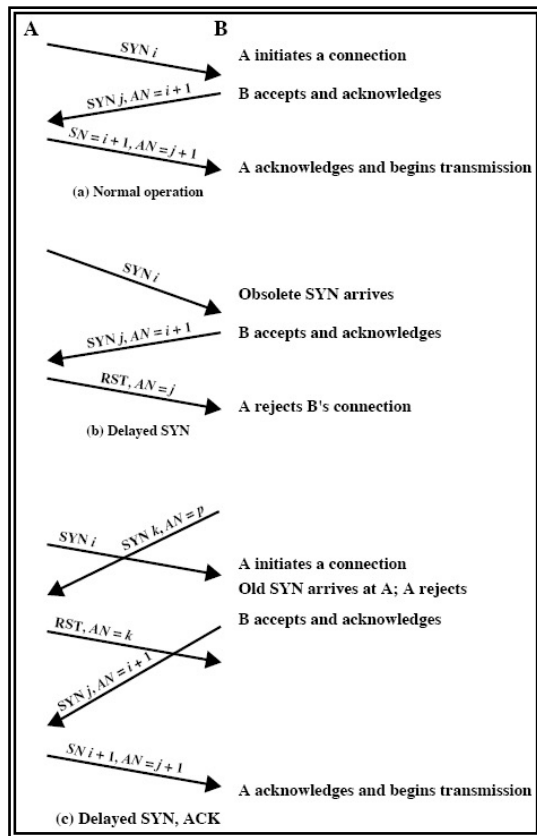
# What if the SYN Packet Gets Lost?

- Suppose the **SYN packet gets lost**
  - Packet dropped by the network or server is busy
- Eventually, **no SYN-ACK arrives**
  - Sender retransmits the SYN on timeout
- How should the **TCP sender set the timer?**
  - Sender has no idea how far away the receiver is
  - Hard to guess a reasonable length of time to wait
  - SHOULD (RFCs 1122 & 2988) use **default of 3 seconds**
    - ✓ Some implementations instead use 6 seconds





# Three-Way Handshake: Examples





# Outline

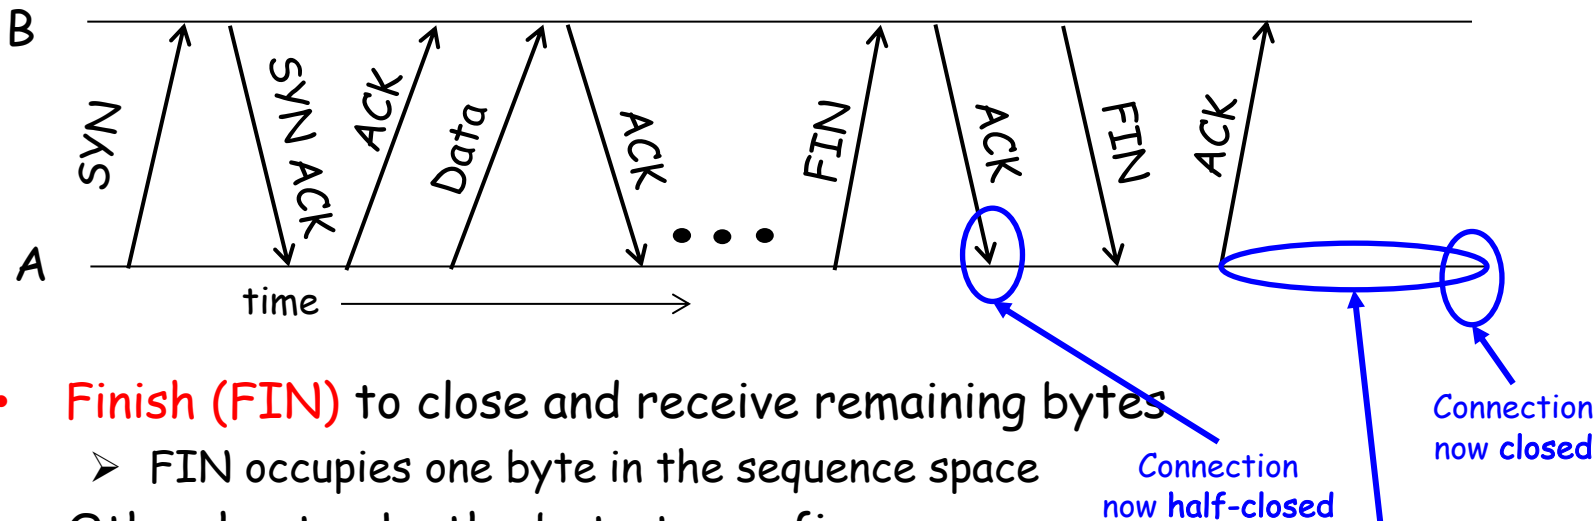
---

- UDP: User Datagram Protocol
- TCP: Transmission Control Protocol
- TCP Connection Setup
- TCP Connection Teardown





# Normal termination, one side at a time



- **Finish (FIN)** to close and receive remaining bytes
  - FIN occupies one byte in the sequence space
- Other host acks the byte to confirm
- Closes A's side of the connection, but not B's
  - Until **B likewise sends a FIN**
  - Which A then acks

**TIME\_WAIT:**

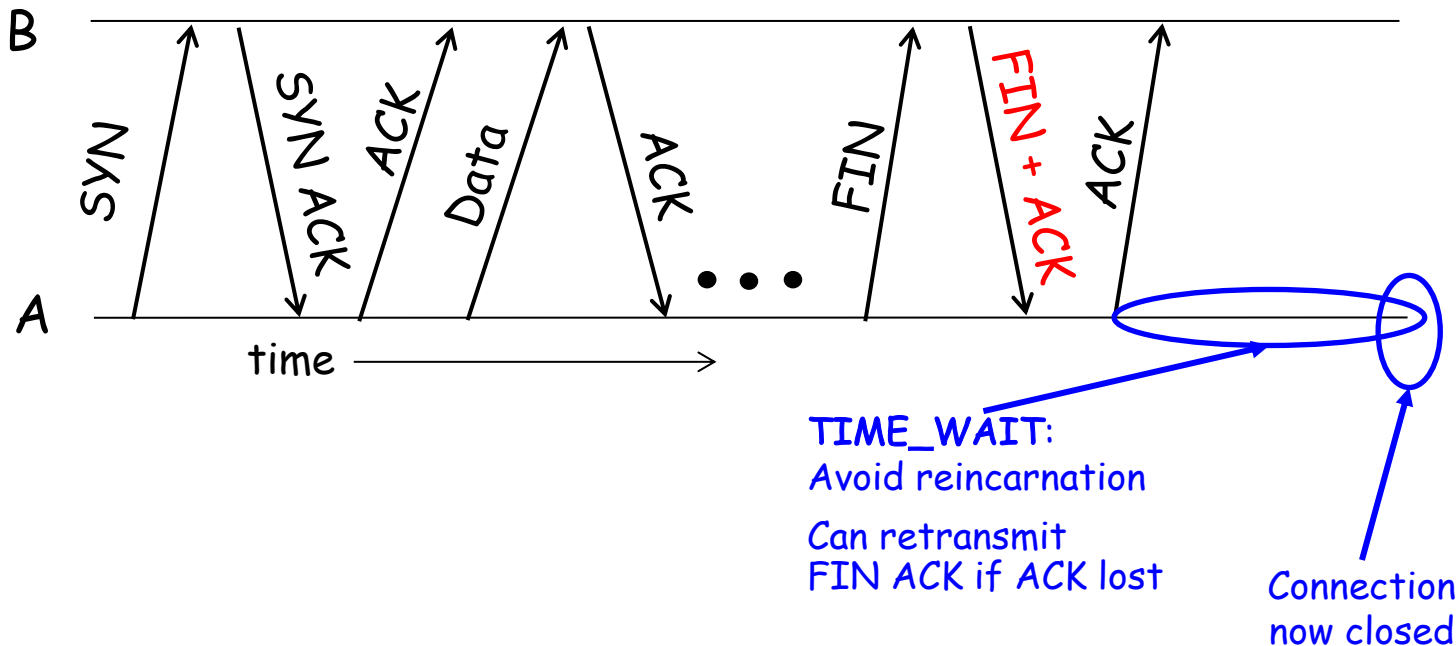
Avoid reincarnation

B will retransmit FIN  
if ACK is lost





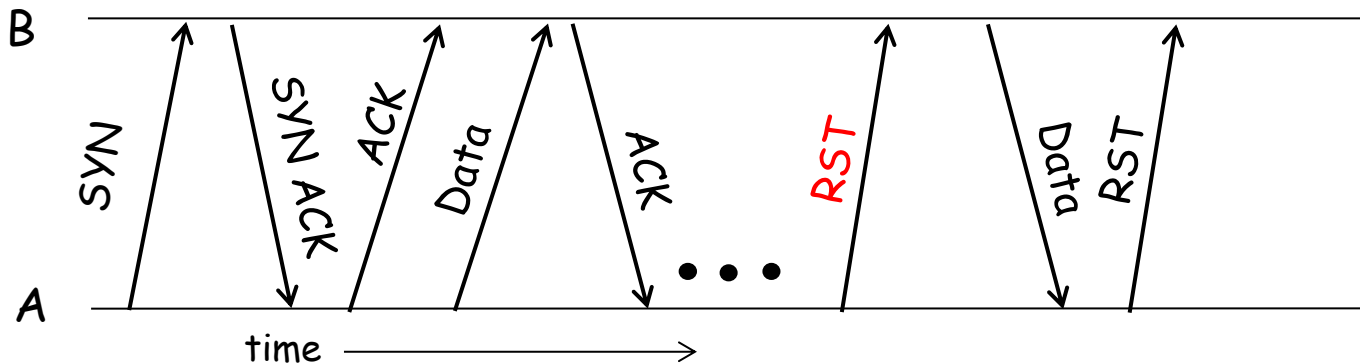
# Normal termination, both together



- Same as before, but **B sets FIN with their ack of A's FIN**



# Abrupt termination

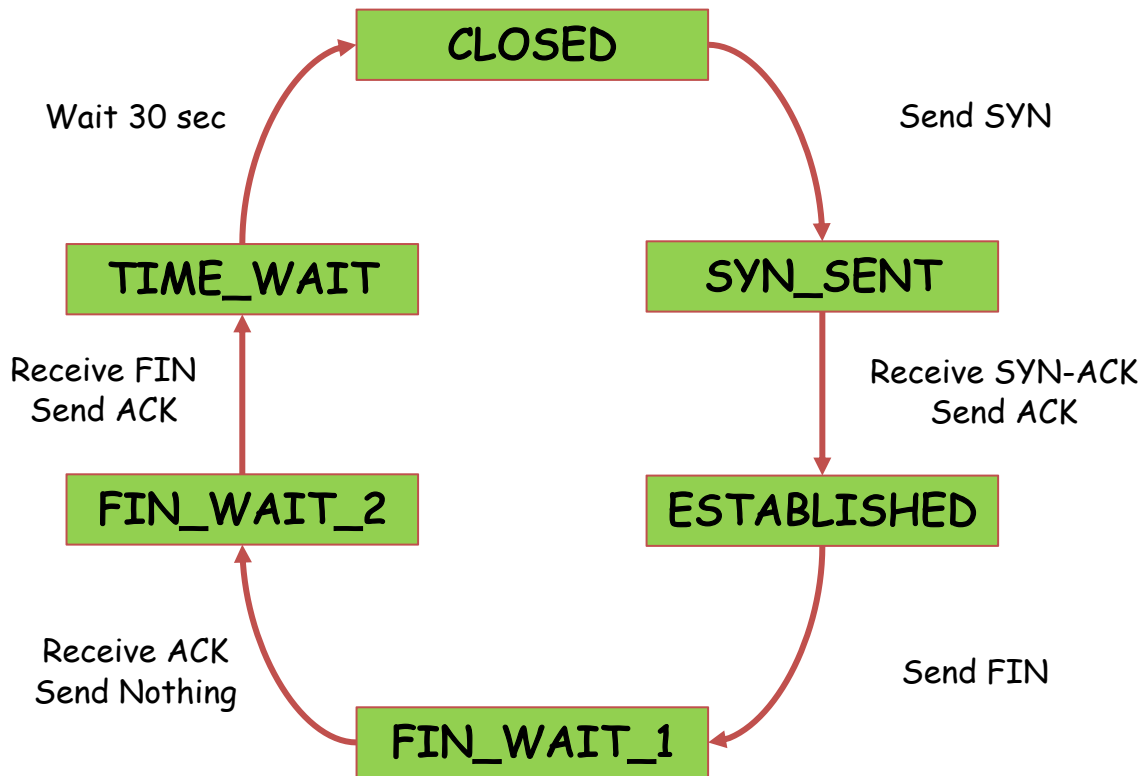


- A sends a **RESET (RST)** to B
  - E.g., because application process on A crashed
- That's it
  - **B does not ack the RST**
  - Thus, RST is not delivered reliably, and any data in flight is lost
  - But: if **B sends anything more, will elicit another RST**



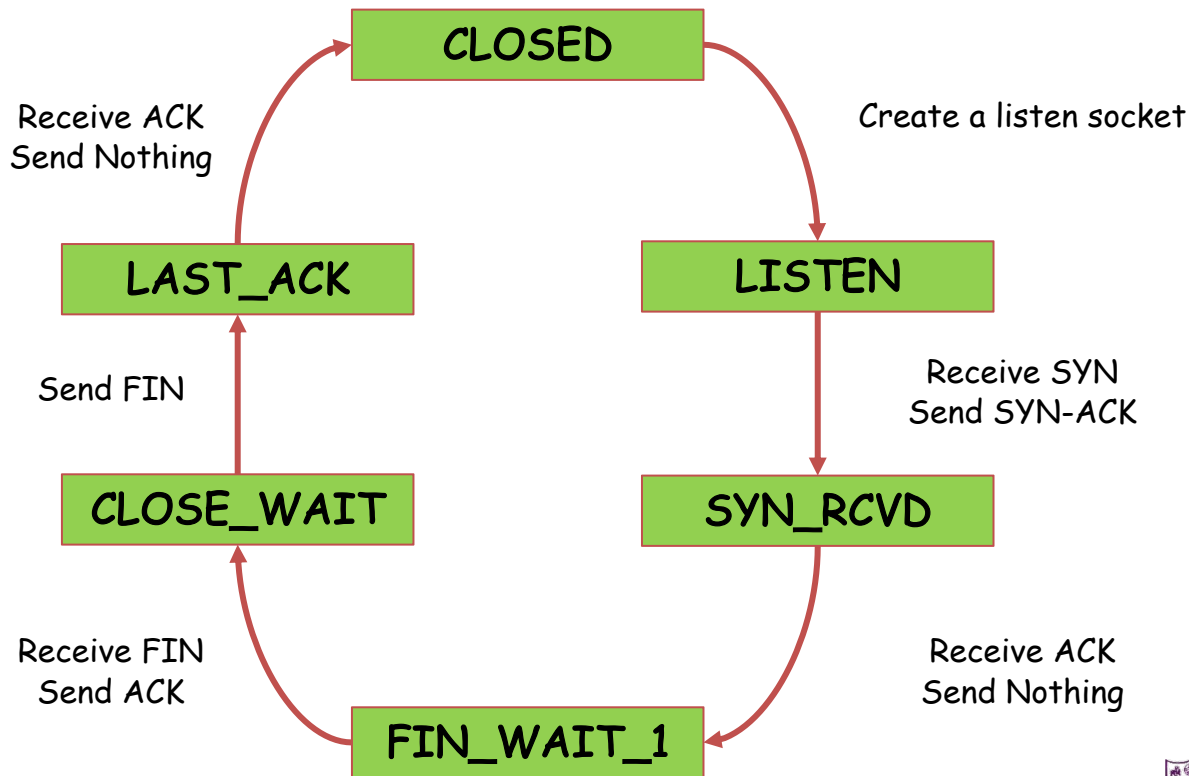


# TCP client lifecycle





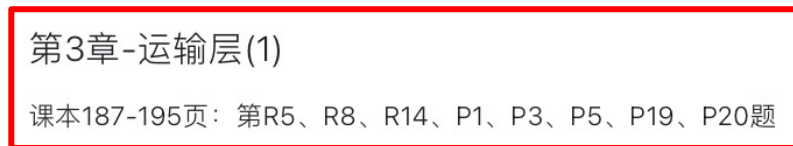
# TCP server lifecycle





# 课程习题（作业）——截止日期：4月8日晚23:59

- **课本187-195页**：第R5、R8、R14、P1、P3、P5、P19、P20题
- 提交方式：<https://selearning.nju.edu.cn/>（教学支持系统）



- 命名：学号+姓名+第\*章。
- 若提交遇到问题请及时发邮件或在下一次上课时反馈。



## 课程习题（作业）——截止日期：4月8日晚23:59

- R5. 在今天的因特网中，为什么语音和图像流量常常是经过 TCP 而不是经 UDP 发送。（提示：我们寻找的答案与 TCP 的拥塞控制机制没有关系。）
- R8. 假定在主机 C 端口 80 上运行的一个 Web 服务器。假定这个 Web 服务器使用持续连接，并且正在接收来自两台不同主机 A 和 B 的请求。被发送的所有请求都通过位于主机 C 的相同套接字吗？如果它们通过不同的套接字传递，这两个套接字都具有端口 80 吗？讨论和解释之。
- R14. 是非判断题：
- a. 主机 A 经过一条 TCP 连接向主机 B 发送一个大文件。假设主机 B 没有数据发往主机 A。因为主机 B 不能随数据捎带确认，所以主机 B 将不向主机 A 发送确认。
  - b. 在连接的整个过程中，TCP 的 `rwnd` 的长度决不会变化。
  - c. 假设主机 A 通过一条 TCP 连接向主机 B 发送一个大文件。主机 A 发送但未被确认的字节数不会超过接收缓存的大小。
  - d. 假设主机 A 通过一条 TCP 连接向主机 B 发送一个大文件。如果对于这条连接的一个报文段的序号为  $m$ ，则对于后继报文段的序号将必然是  $m + 1$ 。
  - e. TCP 报文段在它的首部中有一个 `rwnd` 字段。
  - f. 假定在一条 TCP 连接中最后的 `SampleRTT` 等于 1 秒，那么对于该连接的 `TimeoutInterval` 的当前值必定大于等于 1 秒。
  - g. 假设主机 A 通过一条 TCP 连接向主机 B 发送一个序号为 38 的 4 个字节的报文段。在这个相同的报文段中，确认号必定是 42。



## 课程习题（作业）——截止日期：4月8日晚23:59

- P1. 假设客户 A 向服务器 S 发起一个 Telnet 会话。与此同时，客户 B 也向服务器 S 发起一个 Telnet 会话。给出下面报文段的源端口号和目的端口号：
- 从 A 向 S 发送的报文段。
  - 从 B 向 S 发送的报文段。
  - 从 S 向 A 发送的报文段。
  - 从 S 向 B 发送的报文段。
  - 如果 A 和 B 是不同的主机，那么从 A 向 S 发送的报文段的源端口号是否可能与从 B 向 S 发送的报文段的源端口号相同？
  - 如果它们是同一台主机，情况会怎么样？
- P3. UDP 和 TCP 使用反码来计算它们的检验和。假设你有下面 3 个 8 比特字节：01010011，01100110，01110100。这些 8 比特字节和的反码是多少？（注意到尽管 UDP 和 TCP 使用 16 比特的字来计算检验和，但对于这个问题，你应该考虑 8 比特和。）写出所有工作过程。UDP 为什么要用该和的反码，即为什么不直接使用该和呢？使用该反码方案，接收方如何检测出差错？1 比特的差错将可能检测不出来吗？2 比特的差错呢？







## 课程习题（作业）——截止日期：4月8日晚23:59

- P5. 假定某 UDP 接收方对接收到的 UDP 报文段计算因特网检验和，并发现它与承载在检验和字段中的值相匹配。该接收方能够绝对确信没有出现过比特差错吗？试解释之。
- P19. 考虑一种情况，主机 A 想同时向主机 B 和主机 C 发送分组。A 与 B 和 C 是经过广播信道连接的，即由 A 发送的分组通过该信道传送到 B 和 C。假设连接 A、B 和 C 的这个广播信道具有独立的报文丢失和损坏特性（例如，从 A 发出的报文可能被 B 正确接收，但没有被 C 正确接收）。设计一个类似于停等协议的差错控制协议，用于从 A 可靠地传输分组到 B 和 C。该协议使得 A 直到得知 B 和 C 已经正确接收到当前报文，才获取上层交付的新数据。给出 A 和 C 的 FSM 描述。（提示：B 的 FSM 大体上应当与 C 的相同。）同时，给出所使用的报文格式的描述。
- P20. 考虑一种主机 A 和主机 B 要向主机 C 发送报文的情况。主机 A 和 C 通过一条报文能够丢失和损坏（但不重排序）的信道相连接。主机 B 和 C 由另一条（与连接 A 和 C 的信道独立）具有相同性质的信道连接。在主机 C 上的运输层，在向上层交付来自主机 A 和 B 的报文时应当交替进行（即它应当首先交付来自 A 的分组中的数据，然后是来自 B 的分组中的数据，等等）。设计一个类似于停等协议的差错控制协议，以可靠地向 C 传输来自 A 和 B 的分组，同时以前面描述的方式在 C 处交替地交付。给出 A 和 C 的 FSM 描述。（提示：B 的 FSM 大体上应当与 A 的相同。）同时，给出所使用的报文格式的描述。





# 提问

## Q & A

殷亚凤

智能软件与工程学院

苏州校区南雍楼东区225

yafeng@nju.edu.cn , <https://yafengnju.github.io/>



南京大學  
NANJING UNIVERSITY